

TaxaHabitat

Habitat assignment and spatial quality control for the TaxaID ecosystem. Classifies species into habitat categories using LLM-based biological consensus, assigns habitats to sampling sites, and flags spatial outliers. It accepts that species can occur in multiple habitats, but assumes that a particular sampling location can be defined as a single habitat category. In other words, it uses the species collected at a location to guess the habitat at the site where the species was observed. It gets this right most of the time, and indicates its uncertainty when it can't decide.

Habitat (e.g., terrestrial, marine, freshwater) is a key predictor of which species are plausible at a sampling location. In TaxaID, I recommend a few coarse habitat types rather than a long list of subtle ones. Incorrect habitat classification leads to false positives when species from the wrong habitat receive inflated priors. Ideally, a knowledgeable user assigns each taxon to its habitat. When the species list is long, LLMs can do a reasonable job of classifying species into basic habitat categories. As with any LLM output, users should review the results. `flag_habitat_inconsistencies()` provides an interactive map that makes errant classifications easy to spot and correct.

The interactive map makes it easier to review and proof occupancy data. Even well-curated data like GBIF (Global Biodiversity Information Facility; GBIF Secretariat, Copenhagen, Denmark) have a high frequency of location errors. By mapping points by habitat type, users can easily view which observations have incorrect coordinates. The function `review_spatial_flags(occurrences_flagged)` is designed to flag errant points for removal before model building begins. For instance, a point on the map in the middle of the ocean labeled “Terrestrial”, is likely an error in habitat assignment that would weaken a species distribution model. This tool makes it possible to delete that point or reassign it to “Marine”. TaxaHabitat thus can be a standalone database QAQC for biodiversity databases and helpful for creating species distribution models from occupancy data.

Habitat Schemes

Scheme	Categories	Use case
3-category (default)	Marine / Freshwater / Terrestrial	Most eDNA studies
IUCN Level 1	18 IUCN habitat categories	Fine-grained habitat mapping
Custom	User-defined	Specialized study designs

A custom scheme is a plain data frame (`l1_name` required; `l2_name`, `l2_code`, `realm` optional – see `example_habitat_scheme` for the exact shape):

```
my_scheme <- data.frame(  
  l1_name = c("Kelp forest", "Sandy bottom", "Rocky intertidal"),  
  realm   = c("marine", "marine", "marine")  
)  
prompt <- build_habitat_prompt(taxa, habitat_scheme = my_scheme)
```

Don't want to hand-build one? `build_scheme_prompt()` + `parse_scheme_response()` can draft a scheme from your taxon list via LLM first, and you edit the result before using it.

Installation

```
# Requires TaxaTools (foundation package)  
devtools::install("path/to/TaxaTools")  
devtools::install("path/to/TaxaHabitat")
```

Quick Start

```
library(TaxaHabitat)

# 1. Build an LLM prompt for habitat classification
prompt <- build_habitat_prompt(
  c("Fundulus parvipinnis", "Cottus asper", "Anas platyrhynchos")
)

# 2. Submit to an LLM provider
raw_text <- TaxaTools::prompt_api(prompt)

# 3. Parse response into per-species habitat weights
habitat_weights <- parse_hierarchical_habitat_response(raw_text, prompt)
# taxon_name      Marine_weight  Freshwater_weight  Terrestrial_weight
# Fundulus parvipinnis  0.85            0.15            0.0
# Cottus asper        0.0            1.0            0.0
# Anas platyrhynchos  0.05            0.60            0.35

# Steps 1-3 in one call, with a per-taxon on-disk cache -- what every
# production workflow should use. A habitat verdict decides which occurrence
# records count toward a habitat-stratified site prior downstream, so an
# uncached verdict that differs between two runs moves priors by orders of
# magnitude and makes the published taxon list irreproducible. Only taxa not
# already classified under the same scheme are sent to the LLM.
habitat_weights <- build_habitat_lookup(
  c("Fundulus parvipinnis", "Cottus asper", "Anas platyrhynchos"),
  cache_dir = "my_project_habitat_cache"
)
attr(habitat_weights, "cache_summary") # n_from_cache / n_called

# 4. Assign habitat to sampling sites based on species composition
occurrences_with_habitat <- assign_habitat_biological(
  occurrences, habitat_weights
)

# 5. Flag spatial outliers (marine species at inland sites, etc.)
flagged <- flag_habitat_inconsistencies(occurrences_with_habitat)
```

Key Functions

Habitat classification: - `build_habitat_lookup()` - cached one-call classification (prompt, LLM call, parse; a taxon already classified under the same scheme is never re-asked) - `taxahabitat_clear_cache()` - report or prune that cache - `build_habitat_prompt()` - create LLM prompt for species habitat weights - `parse_hierarchical_habitat_response()` - parse LLM output to numeric weights - `assign_habitat_biological()` - assign site habitat from species composition - `consensus_habitat()` - assemblage-level consensus with ecoregion extraction

Custom schemes: - `build_iucn_scheme()` - generate IUCN Level 1 habitat scheme - `example_habitat_scheme()` - example custom scheme for reference - `build_scheme_prompt()` / `parse_scheme_response()` - custom scheme workflow

Spatial QC: - `flag_habitat_inconsistencies()` – flag records inconsistent with site habitat - `review_spatial_flags()` – interactive Leaflet map for manual review

Reporting: - `report_habitat()` – summarize habitat assignment for `assemble_report()`

Methods

Habitat Weight Estimation

Rather than assign each species to a single habitat, TaxaHabitat asks the LLM to distribute habitat affinity as continuous weights across all habitat categories (e.g., Marine 0.85, Freshwater 0.15, Terrestrial 0.0), summing to 1.0 per species. This captures habitat generalism – an estuarine fish contributes partial signal to both Marine and Freshwater – and avoids the information loss of a categorical assignment.

Prompts are constructed by `build_habitat_prompt()` and sent to an LLM provider via TaxaTools. For large species lists, taxa are chunked (default 60 per call) into self-contained prompts. The LLM returns CSV with numeric weights per habitat plus an `Other_weight` column and free-text `habitat_best_guess` for species that do not fit the scheme. `parse_hierarchical_habitat_response()` validates that row sums are within tolerance (warns if deviation > 0.05 from 1.0) and folds unrecognized habitat columns into `Other_weight`.

Site Habitat Assignment

`assign_habitat_biological()` assigns a habitat to each sampling location based on the species observed there. For each point, the function joins occurrence records to species-level habitat weights, sums weight vectors across species, normalizes to proportions, and assigns the habitat with the highest proportion if it exceeds a threshold (default 0.3). The threshold is lower than a simple majority (0.5) because generalist species spread weight across multiple habitats, diluting any single category.

By default, each species counts equally regardless of how many times it was recorded at a point (`weight_by_abundance = FALSE`). This is because I assume occurrence record counts reflect sampling effort more than habitat.

Assemblage-Level Consensus

`consensus_habitat()` infers site habitat from a species list alone (without occurrence coordinates), using the same weighted-voting logic. When an `ecoregion_best_guess` column is present (from LLM output with `geographic_context`), the modal ecoregion across species is returned alongside the habitat consensus.

Spatial Quality Control

TaxaHabitat can do some basic ground truthing of its habitat assignments using habitat polygons to overlay the search area. `flag_habitat_inconsistencies()` checks whether each occurrence record is spatially consistent with its assigned habitat using vector polygons (Natural Earth land/ocean boundaries; a public-domain map dataset maintained by volunteer cartographers with support from the North American Cartographic Information Society, naturalearthdata.com) and raster bathymetry (GEBCO, the General Bathymetric Chart of the Oceans; maintained by the GEBCO Compilation Group and hosted by the British Oceanographic Data Centre, Liverpool, United Kingdom). Each record is classified into a physical zone:

Zone	Definition
inland	Inside land polygon, > 1 km from coast
coastal	Within 1 km of coastline (either side)
marine_shallow	Ocean, 0–200 m depth (continental shelf)

Zone	Definition
marine_deep	Ocean, 200–4000 m depth (bathyal)
marine_abyssal	Ocean, > 4000 m depth

Records are then cross-referenced against the habitat’s expected realm (marine, freshwater, terrestrial) and flagged at three levels: “likely” (consistent), “questionable” (borderline, e.g., marine species very near shore), or “unlikely” (clear mismatch, e.g., terrestrial species in open ocean). The 1 km coastal buffer accounts for GPS uncertainty and tidal gradients. `review_spatial_flags()` provides an interactive Leaflet map for manual inspection and correction.

The Natural Earth and GEBCO reference layers themselves are fixed – `flag_habitat_inconsistencies()` does not currently accept a user-supplied coastline, bathymetry, or other custom spatial reference layer. Only the numeric thresholds above (`coast_buffer_m`, `depth_neritic_m`, `depth_oceanic_m`, `bathymetry_resolution`) are tunable.

`review_spatial_flags()` (real Great Lakes data shown above) is where `flag_habitat_inconsistencies()`’s output actually gets reviewed: points are colored by habitat category, clicking one shows its species list, and the sidebar’s Flag / Reassign Habitat controls let you correct a misclassified point without leaving the map.

LLM Integration

TaxaHabitat uses the `llm_fn` pattern from TaxaTools. The default provider is `call_anthropic_api()` (Anthropic Claude; Anthropic PBC, San Francisco, California), but any compatible provider works:

```
# Use Gemini instead
raw_text <- TaxaTools::prompt_api(prompt, llm_fn = TaxaTools::call_gemini_api)

# Same, through the cached wrapper (pass a new cache_tag, or clear the
# cache, when you change model and want fresh verdicts)
habitat_weights <- build_habitat_lookup(taxa, llm_fn = TaxaTools::call_gemini_api,
                                       cache_dir = "my_project_habitat_cache",
                                       cache_tag = "gemini")
```

`call_gemini_api()` calls Google Gemini (Google LLC, Mountain View, California); other supported providers include OpenAI (OpenAI OpCo, LLC, San Francisco, California) and local Ollama (Ollama, Palo Alto, California) – see TaxaTools.

API Keys

Requires an Application Programming Interface (API) key for at least one LLM provider (Anthropic by default). See the TaxaTools API Setup vignette for configuration.

Vignettes

- Habitat Assignment – full workflow guide

Part of TaxaID

TaxaHabitat receives occurrence data from TaxaFetch and produces habitat-annotated records for TaxaExpect (prior estimation).

Ecosystem: TaxaTools -> TaxaFetch -> **TaxaHabitat** -> TaxaExpect -> TaxaAssign

See the TaxaID README for ecosystem overview and installation instructions.

Citation

Lafferty, K.D., 2026, TaxaID – A modular R ecosystem for Bayesian taxonomic assignment: U.S. Geological Survey software release, <https://doi.org/10.5066/xxxxxx>.

Software Requirements

- R ($\geq 4.1.0$; R Core Team 2025)
- TaxaTools (for LLM provider functions)
- An LLM API key is required for habitat assignment via `build_habitat_prompt()`

All dependencies are declared in the DESCRIPTION file and installed automatically.

Developed with Claude Code (Anthropic PBC, San Francisco, California).

References

R Core Team (2025). R: A Language and Environment for Statistical Computing. V.4.5.2. R Foundation for Statistical Computing, Vienna, Austria. <https://www.r-project.org>